

weft 1.1.0

A Modern Literate Programming Tool

User Guide

Cuauhtemoc Pacheco

Derived from nuweb 1.64 by Preston Briggs¹
with contributions by Keith Harwood, Simon Wright,
Marc W. Mengel, John D. Ramsdell, and Gregor Goldbach

March 5, 2026

¹Original nuweb author.

Contents

1	Introduction	1
1.1	weft	1
1.2	Writing weft	3
1.2.1	Source File Extension: <code>.weft</code>	3
1.2.2	The Major Commands	4
1.2.3	The Minor Commands	7
1.3	Sectioning commands	8
1.4	Running weft	9
1.5	Markdown Output	10
1.6	Section Markers and Language Auto-Detection	11
1.6.1	Overriding the Language with <code>@1</code>	11
1.6.2	Mid-Line Inline Expansion	12
1.6.3	The <code>@L</code> Command: Defining New Languages	13
1.6.4	The Language Table	13
1.6.5	Syntax Highlighting in Markdown Output	14
1.6.6	Output Format of Section Markers	14
1.6.7	Debugging Tangled Code	14
1.7	Restrictions	18
1.8	Project Introspection	19
1.8.1	Map (<code>-m</code>)	19
1.8.2	Extract (<code>-e</code>)	19
1.8.3	Reverse Map (<code>-R</code>)	19
1.9	Acknowledgements	20

Chapter 1

Introduction

In 1984, Knuth introduced the idea of *literate programming* and described a pair of tools to support the practise [1]. His approach was to combine Pascal code with T_EX documentation to produce a new language, WEB, that offered programmers a superior approach to programming. He wrote several programs in WEB, including `weave` and `tangle`, the programs used to support literate programming. The idea was that a programmer wrote one document, the web file, that combined documentation (written in T_EX [2]) with code (written in Pascal).

Running `tangle` on the web file would produce a complete Pascal program, ready for compilation by an ordinary Pascal compiler. The primary function of `tangle` is to allow the programmer to present elements of the program in any desired order, regardless of the restrictions imposed by the programming language. Thus, the programmer is free to present his program in a top-down fashion, bottom-up fashion, or whatever seems best in terms of promoting understanding and maintenance.

Running `weave` on the web file would produce a T_EX file, ready to be processed by T_EX. The resulting document included a variety of automatically generated indices and cross-references that made it much easier to navigate the code. Additionally, all of the code sections were automatically prettyprinted, resulting in a quite impressive document.

Knuth also wrote the programs for T_EX and METAFONT entirely in WEB, eventually publishing them in book form [3, 4]. These are probably the largest programs ever published in a readable form.

Inspired by Knuth's example, many people have experimented with WEB. Some people have even built web-like tools for their own favorite combinations of programming language and typesetting language. For example, CWEB, Knuth's current system of choice, works with a combination of C (or C++) and T_EX [6]. Another system, FunnelWeb, is independent of any programming language and only mildly dependent on T_EX [8]. Inspired by the versatility of FunnelWeb and by the daunting size of its documentation, Preston Briggs decided to write his own, very simple, tool for literate programming.¹

`weft` is a fork of `nuweb`. The original `nuweb` code, written by Preston Briggs and maintained over the years by several contributors, provided an excellent foundation: a simple, language-independent literate programming tool. `weft` builds on that foundation with a modernized codebase, a modular source structure organized by concept, the removal of the unmaintained HTML output path, and the addition of a Markdown weave output format. The goal is to continue evolving the tool while preserving the simplicity and philosophy that made `nuweb` valuable in the first place.

1.1 `weft`

`weft` works with any programming language and produces documentation in L^AT_EX [5] or Markdown format. L^AT_EX was the original choice because it supports a multi-level sectioning scheme and has facilities for drawing figures. Markdown support was added to allow literate programs to produce lightweight documentation

¹There is another system similar to `nuweb`, written by Norman Ramsey, called `noweb` [7]. It perhaps suffers from being overly Unix-dependent and requiring several programs to use. On the other hand, its command syntax is very nice. In any case, `nuweb` certainly owes its name and a number of features to his inspiration.

that renders well on platforms like GitHub, GitLab, and any Markdown viewer, without requiring a $\text{\TeX}$ installation.

Support for arbitrary programming languages was a design goal because programmers work in many languages (and sometimes combinations of several languages), *e.g.*, C, Fortran, C++, yacc, lex, Scheme, assembly, Postscript, and so forth. The need to support arbitrary programming languages has many consequences:

No prettyprinting Both `WEB` and `CWEB` are able to prettyprint the code sections of their documents because they understand the language well enough to parse it. Since `weft` aims to support *any* language, this feature must be abandoned. However, `weft` does allow particular individual formulas or fragments of $\text{\LaTeX}$ code to be formatted and still be parts of output files. Also, keywords in scraps can be surrounded by `@_` to have them be bold in the output.

No index of identifiers Because `WEB` knows about Pascal, it is able to construct an index of all the identifiers occurring in the code sections (filtering out keywords and the standard type identifiers). Unfortunately, this isn't as easy for a language-independent tool. `weft` doesn't know what an identifier looks like in each language and certainly doesn't know all the keywords. (On the other hand, see the end of Section 1.2.3)

Of course, there must be some compensation for these losses or the whole idea would be a waste. Here are the advantages:

Simplicity The majority of the commands in `WEB` are concerned with control of the automatic prettyprinting. Since `weft` doesn't prettyprint, many commands are eliminated. A further set of commands is subsumed by $\text{\LaTeX}$ and may also be eliminated. As a result, the set of commands is reduced to only four members (explained in the next section). This simplicity is also reflected in the size of this tool, which is quite a bit smaller than the tools used with other approaches.

No prettyprinting Everyone disagrees about how their code should look, so automatic formatting annoys many people. One approach is to provide ways to control the formatting. `weft`'s approach is simpler—it performs no automatic formatting and therefore allows the programmer complete control of code layout. `weft` does allow individual scraps to be presented in either verbatim, math, or paragraph mode in the $\text{\TeX}$ output.

Control `weft` also offers the programmer complete control of the layout of his output files (the files generated during tangling). Of course, this is essential for languages that are sensitive to layout; but it is also important in many practical situations, *e.g.*, debugging.

Speed Since `weft` doesn't do too much, the `weft` tool runs quickly. By default it performs only tangling; weaving is activated with `-w`. When weaving is requested, both operations execute in a single pass.

Page numbers Inspired by the example of `noweb`, `weft` refers to all scraps by page number to simplify navigation. If there are multiple scraps on a page (say, page 17), they are distinguished by lower-case letters (*e.g.*, 17a, 17b, and so forth).

Multiple file output The programmer may specify more than one output file in a single `weft` file. This is required when constructing programs in a combination of languages (say, Fortran and C). It's also an advantage when constructing very large programs that would require a lot of compile time.

This last point is very important. By allowing the creation of multiple output files, the need for monolithic programs is avoided. This supports the creation of very large programs by groups of people.

A further reduction in compilation time is achieved by first writing each output file to a temporary location, then comparing the temporary file with the old version of the file. If there is no difference, the temporary file can be deleted. If the files differ, the old version is deleted and the temporary file renamed. This approach works well in combination with `make` (or similar tools), since `make` will avoid recompiling untouched output files.

1.2 Writing weft

The bulk of a weft file will be ordinary L^AT_EX (or Markdown, when using the `-w md` flag). In fact, any L^AT_EX file can serve as input to weft and will be simply copied through, unchanged, to the documentation file—unless a weft command is discovered. Similarly, any Markdown file can be processed with `-w md` and will be copied through to the `.md` output. All weft commands begin with an “at-sign” (`@`). Therefore, a file without at-signs will be copied unchanged. weft commands are used to specify *output files*, define *fragments*, and delimit *scraps*. These are the basic features of interest to the weft tool—all else is simply text to be copied to the documentation file.

1.2.1 Source File Extension: `.weft`

weft source files use the `.weft` extension. This single, format-agnostic extension replaces the earlier `.w` and `.wmd` conventions. The prose surrounding scraps may be written in L^AT_EX or Markdown; the at-sign commands (`@o`, `@d`, `@<`, etc.) are the same in both cases.

The weave output format is specified in one of two ways:

In the source file The `@W` directive declares the intended format: `@W tex` for L^AT_EX, `@W md` for Markdown. This should appear in the preamble, before any `@o` or `@d` directives.

On the command line `-w md` or `-w tex` overrides the `@W` directive. Using `-w` alone (without a format token) reads the format from `@W`.

If no extension is given on the command line, weft appends `.weft` automatically.

Default Behavior: Tangle Only

By default, `weft myfile.weft` performs only tangling (producing the output files). No documentation is generated unless `-w` is given:

<code>weft programa.weft</code>	tangle only (default)
<code>weft -w programa.weft</code>	weave per <code>@W</code> (error if missing)
<code>weft -w md programa.weft</code>	weave Markdown (override)
<code>weft -w tex programa.weft</code>	weave LaTeX (override)

The `-t` flag suppresses weaving even if `-w` is given. The `-o` flag suppresses tangling.

Cosmetic Fences

When writing `.weft` files with Markdown prose, it is often convenient to surround scrap definitions with Markdown fenced code blocks (triple backticks) so that the source file renders nicely in Markdown previewers such as GitHub or VS Code. For example:

Some prose in Markdown.

```
'''  
@d My fragment  
@{code here  
@}  
'''
```

More prose.

weft silently strips these fence lines (lines that consist of three or more backticks, optionally followed by a language identifier) during Markdown weaving. They do not appear in the generated `.md` file and do not affect tangling. This is purely a cosmetic convenience for making the source file readable in any Markdown viewer.

Note that fences are only stripped in Markdown mode (`-w md`). In L^AT_EX mode, backtick lines are treated as ordinary text.

1.2.2 The Major Commands

Files and fragments are defined with the following commands:

@o *file-name flags scrap* Output a file. The file name is terminated by whitespace.

@d *fragment-name scrap* Define a fragment. The fragment name is terminated by a return or the beginning of a scrap.

@q *fragment-name scrap* Define a fragment. The fragment name is terminated by a return or the beginning of a scrap. This a quoted fragment.

A specific file may be specified several times, with each definition being written out, one after the other, in the order they appear. The definitions of fragments may be similarly specified piecemeal.

Fragment Names with Embedded Parameters

Fragment names can carry *parameters*—placeholder values embedded directly in the name using the **@'value@'** syntax. Inside the fragment body, **@1**, **@2**, ... **@9** refer to the first, second, ... ninth parameter in order of appearance.

For example, a reusable validation fragment:

```
@d Validate field @'name@' type @'string@'
@{if (!data.@1) throw new Error("Missing @1");
if (typeof data.@1 !== '@2')
    throw new TypeError("@1 must be @2");
@}
```

Each invocation supplies its own arguments:

```
<@Validate field @'email@' type @'string@'>
<@Validate field @'age@' type @'number@'>
```

The first invocation expands to:

```
if (!data.email) throw new Error("Missing email");
if (typeof data.email !== 'string')
    throw new TypeError("email must be string");
```

Note three key points:

- The values in the *definition* (**@'name@'**, **@'string@'**) serve as default arguments and as documentation—they make the fragment name readable.
- At each *invocation*, different values can be substituted (**@'email@'** replaces **@1**, **@'string@'** replaces **@2**).
- If a fragment name is abbreviated with ..., any missing arguments fall back to the defaults given in the definition.

The difference between a quoted fragment (**@q**) and an ordinary one (**@d**) is that inside a quoted fragment fragments are not expanded on output. Rather, they are formatted as uses of fragments so that the output file can itself be **weft** source. This allows you to create files containing fragments which can undergo further processing before the fragments are expanded, while also describing and documenting them in one place.

You can have both quoted and unquoted fragments with the same name. They are written out in order as usual, with those introduced by **@q** being quoted and those with **@d** expanded as normal.

In quoted fragments, the **@f** filename is quoted as well, so that when it is expanded it refers to the finally produced file, not any of the intermediate ones.

Scraps

Scraps have specific begin markers and end markers to allow precise control over the contents and layout. Note that any amount of whitespace (including carriage returns) may appear between a name and the beginning of a scrap. Scraps may also appear in the running text where they are formatted (almost) identically to their use in definitions, but don't appear in any code output.

`@{anything@}` where the scrap body includes every character in *anything*—all the blanks, all the tabs, all the carriage returns. This scrap will be typeset in verbatim mode. Using the `-l` option will cause the program to typeset the scrap with the help of L^AT_EX's `listings` package.

`@[anything@]` where the scrap body includes every character in *anything*—all the blanks, all the tabs, all the carriage returns. This scrap will be typeset in paragraph mode, allowing sections of T_EX documents to be scraps, but still be prettyprinted in the document.

`@(anything@)` where the scrap body includes every character in *anything*—all the blanks, all the tabs, all the carriage returns. This scrap will be typeset in math mode. This allows this scrap to contain a formula which will be typeset nicely.

Inside a scrap, several special sequences are recognized.

Fragment invocation.

`<fragment-name>` Expands the named fragment inline. It is an error to specify recursive invocations. If the fragment name contains embedded parameters, the invocation supplies arguments at the same positions (see above).

When the reference appears *mid-line*—that is, after at least one visible character on the same line—section markers, `#line` directives, and fragment-name comments are automatically suppressed. The fragment content is spliced seamlessly into the surrounding text. At start-of-line positions (including indented lines), markers are emitted normally. See Section 1.6 for details.

Argument substitution.

`@1`, `@2`, ..., `@9` Inside a fragment body, substitutes the *n*th argument. If the argument was not supplied, the empty string is used.

At each argument position in an invocation, four forms are accepted:

`@'string@'` A literal string, copied verbatim (only `@@` is interpreted, producing a single `@`). This is the most common form.

`<fragment-name>` The fragment is expanded and its output becomes the argument value.

`@{text@}` The text is expanded as an anonymous fragment (inheriting the caller's arguments) and the result is passed as the argument. Useful for combining text and arguments into a single value.

`@1`, `@2`, ..., `@9` Forwards an argument from the calling fragment to the called one.

Other in-scrap commands.

`@xlabel@x` Marks a place inside a scrap and associates it to the label (which can be any text not containing a `@`). Expands to the reference number of the scrap followed by a numeric value. Outside scraps it expands to the same value. It is used so that text outside scraps can refer to particular places within scraps.

`@f` Inside a scrap this is replaced by the name of the current output file.

`@t` Inside a scrap this is replaced by the title of the fragment as it is at the point it is used, with all parameters replaced by actual arguments.

@# At the beginning of a line in a scrap this will suppress the normal indentation for that line. Use this, for example, when you have a `#ifdef` inside a nested scrap. Writing `@##ifdef` will cause it to be lined up on the left rather than indented with the rest of its code.

@s Inside a scrap suppresses indentation for the following fragment expansion.

Note that fragment names may be abbreviated, either during invocation or definition. For example, it would be very tedious to have to type, repeatedly, the fragment name

```
@d Check for terminating at-sequence and return name if found
```

Therefore, `weft` provides a mechanism (borrowed from Knuth) of indicating abbreviated names.

```
@d Check for terminating...
```

Basically, the programmer need only type enough characters to identify the fragment name uniquely, followed by three periods. An abbreviation may even occur before the full version; `weft` simply preserves the longest version of a fragment name. Note also that blanks and tabs are insignificant within a fragment name; each string of them is replaced by a single blank.

Sometimes, for instance during program testing, it is convenient to comment out a few lines of code. In C or Fortran placing `/* ... */` around the relevant code is not a robust solution, as the code itself may contain comments. `weft` provides the command

```
@%
```

only to be used inside scraps. It behaves exactly the same as `%` in the normal L^AT_EX text body.

When scraps are written to a program file or a documentation file, tabs are expanded into spaces by default. Currently, `weft` assumes tab stops are set every eight characters. Furthermore, when a fragment is expanded in a scrap, the body of the fragment is indented to match the indentation of the fragment invocation. Therefore, care must be taken with languages (*e.g.*, Fortran) that are sensitive to indentation. These default behaviors may be changed for each output file (see below).

Flags

When defining an output file, the programmer has the option of using flags to control output of a particular file. The flags are intended to make life a little easier for programmers using certain languages. They introduce little language dependences; however, they do so only for a particular file. Thus it is still easy to mix languages within a single document. There are five “per-file” flags:

- d Forces the creation of `#line` directives in the output file. These are useful with C (and sometimes C++ and Fortran) on many Unix systems since they cause the compiler’s error messages to refer to the web file rather than to the output file. Similarly, they allow source debugging in terms of the web file.
- i Suppresses the indentation of fragments. That is, when a fragment is expanded within a scrap, it will *not* be indented to match the indentation of the fragment invocation. This flag would seem most useful for Fortran programmers.
- t Suppresses expansion of tabs in the output file. This feature seems important when generating `make` files.
- cx Puts comments in generated code documenting the fragment that generated that code. The *x* selects the comment style for the language in use: `c` for C (`/* */`), `+` for C++ (`//`), and `p` for shell/Perl (`#`). If the global `-x` cross-reference flag is set the comment includes the page reference for the first scrap that generated the code.
- s Suppresses section markers and `#line` directives in the tangled output for this file. The language is still detected (from the extension or a preceding `@1`), so the Markdown weaver can emit syntax-highlighting classes, but the tangled code is emitted without any tangle annotations. This is useful for bootstrap source files or generated output that should remain clean.

These flags are still available for explicit control, but in most cases the automatic language detection described in Section 1.6 is more convenient.

1.2.3 The Minor Commands

There are some very low-level utility commands that may appear anywhere in the web file.

@@ Causes a single “at sign” to be copied into the output.

@_ Causes the text between it and the next **@_** to be made bold (for keywords, etc.)

@i file-name Includes a file. Includes may be nested, though there is currently a limit of 10 levels. The file name should be complete (no extension will be appended) and should be terminated by a carriage return. Normally the current directory is searched for the file to be included, but this can be varied using the **-I** flag on the command line. Each such flag adds one directory to the search path and they are searched in the order given.

@rx Changes the escape character ‘@’ to ‘x’. This must appear before any scrap definitions.

@v Always replaced by the string established by the **-V** flag, or a default string if the flag isn’t given. This is intended to mark versions of the generated files.

@l tag Sets the language for the next **@o** or **@d** directive. The tag is looked up in the language table (Section 1.6) to determine the comment style for section markers and the **class** attribute for Markdown syntax highlighting. Placed on the line before the directive it applies to.

@L name style [+d] Defines a custom language mapping for the section marker system (see Section 1.6). This must appear in the document preamble, before any **@o** or **@d** directives that depend on it.

@W format Declares the weave output format for this source file. Valid formats are **md** (Markdown) and **tex** (L^AT_EX). Used by **-w** when no explicit format is given on the command line. A CLI format (**-w md** or **-w tex**) overrides this directive. Should appear in the document preamble.

In the text of the document, that is outside scraps, you may include scrap-like material.

@{Anything@} The included material, the *Anything*, is typeset as if it appeared inside a scrap. This is useful for referring to fragments in the text and for describing the literate programming process itself.

@<Fragment name@> The fragment named is expanded in place in the text. The expansion is presented verbatim, it is not interpreted for typesetting, so any special environment must be set up before and after this is used.

@c Outside of a scrap this is used to define a block comment using all the text until the next command. Inside a scrap the last block comment is expanded using the current indentation and breaking lines between words if 60 characters are exceeded.

Finally, there are three commands used to create indices to the fragment names, file definitions, and user-specified identifiers.

@f Create an index of file names.

@m Create an index of fragment names.

@u Create an index of user-specified identifiers.

These are usually placed in their own section in the L^AT_EX document; for example, see Chapter ??.

Identifiers must be explicitly specified for inclusion in the **@u** index. By convention, each identifier is marked at the point of its definition; all references to each identifier (inside scraps) will be discovered automatically. To “mark” an identifier for inclusion in the index, it must be mentioned at the end of a scrap. For example,

```
@d a scrap @{
  Suppose the variables FOO and BAR are being declared
  inside this scrap.
@| FOO BAR @}

```

Alphabetic identifiers have been used in this example, but any string of characters (not including whitespace or @ characters) will do. Therefore, it's possible to add index entries for things like <=< if desired. An identifier may be declared in more than one scrap.

In the generated index, each identifier appears with a list of all the scraps using and defining it, where the defining scraps are distinguished by underlining. Note that the identifier doesn't actually have to appear in the defining scrap; it just has to be in the list of definitions at the end of a scrap.

1.3 Sectioning commands

For larger documents the indexes and usage lists get rather unwieldy and problems arise in naming things so that different things in different parts of the document don't get confused. weft has a sectioning command which keeps the fragment names and user identifiers separate. Thus, you can give a fragment in one section the same name as a fragment in another and the two won't be confused or connected in any way. Nor will user identifiers defined in one section be referenced in another. Except for the fact that scraps in successive sections can go into the same output file, this is the same as if the sections came from separate input files.

However, occasionally you may really want fragments from one section to be used in another. More often, you will want to identify a user identifier in one section with the same identifier in another (as, for example, a header file defined in one section is included in code in another). Extra commands allow a fragment defined in one section to be accessible from all other sections. Similarly, you can have scraps which define user identifiers and export them so that they can be used in other sections.

@s Start a new section.

@S Close the current section and don't start another.

@d+ **fragment-name scrap** Define a fragment which is accessible in all sections, a global fragment.

@D+ Likewise

@q+ Likewise

@Q+ Likewise

@m+ Create an index of all such fragments.

@u+ Create an index of globally accessible user identifiers.

There are two kinds of section, the base section which is where normally everything goes, and local sections which are introduced with the @s command. A local section comprises everything from the command which starts it to the one which ends it. A @s command will start a new local section. A @S command closes the current local section, but doesn't open another, so what follows goes into the base section. Note that fragments defined in the base section aren't global; they are accessible only in the base section, but they are accessible regardless of any local sections between their definition and their use.

Within a scrap:

@<+*fragment-name*@> Expand the globally accessible fragment with that name, rather than any local fragment.

@+ Like @| except that the identifiers defined are exported to the global realm and are not directly referenced in any scrap in any section (not even the one where they are defined).

@- Like @| except that the identifiers are imported to the local realm. The cross-references show where the global variables are defined and defines the same names as locally accessible. Uses of the names within the section will point to this scrap.

Note that the + signs above are part of the commands. They are not part of the fragment names. If you want a fragment whose name begins with a plus sign, leave a space between the command and the name.

1.4 Running weft

weft is invoked using the following command:

```
weft flags file-name...
```

One or more files may be processed at a time. If a file name has no extension, `.weft` will be appended. While a file name may specify a file in another directory, the resulting documentation file will always be created in the current directory. For example,

```
weft /foo/bar/quux
```

will take as input the file `/foo/bar/quux.weft` and will produce the tangled output files only. To also generate documentation, use `-w` (see below).

By default, weft performs only tangling—it generates the output files defined by `@o` directives and nothing else. Documentation (weaving) is produced only when the `-w` flag is given. When `-w` is present, both tangling and weaving execute in a single pass. Three flags control what weft produces:

- t Suppress weaving (useful to override `-w` in scripts).
- o Suppress generation of the output files (weave only).
- c Avoid testing output files for change before updating them.

Thus, the command

```
weft -to /foo/bar/quux
```

would simply scan the input and produce no output at all.

There are several additional command-line flags:

- v For “verbose”, causes weft to write information about its progress to `stderr`.
- n Forces scraps to be numbered sequentially from 1 (instead of using page numbers). This form is perhaps more desirable for small webs.
- s Doesn’t print list of scraps making up each file following each scrap.
- d Print “dangling” identifiers – user identifiers which are never referenced, in indices, etc.
- p *path* Prepend *path* to the filenames for all the output files.
- w [*format*] Activate weaving (documentation generation). Without `-w`, no documentation file is produced. The optional *format* argument selects the output format:
 - `-w md` — force Markdown output (`.md` file).
 - `-w tex` — force L^AT_EX output (`.tex` file).
 - `-w` alone — read the format from the `@W` directive in the source file. If `@W` is absent, weft reports an error.

A format given on the command line overrides any `@W` directive. When Markdown output is selected, sequential numbering (`-n`) is forced automatically, since Markdown has no concept of page numbers. The Markdown output uses HTML anchors (`<a id="...">`) for cross-references and `<pre><code>` blocks for scraps, allowing clickable navigation between definitions and references. Cosmetic fence lines (triple backticks) in the source are silently stripped.

- `-x` Include cross-reference numbers in scrap comments.
- `-h options` Provide options for the hyperref package.
- `-r` Turn on hyperlinks. You must include the `—usepackage—` options in the text.
- `-l` Use the `listings` package for formatting scraps. Use this if you want to have a pretty-printer for your scraps. In order to e.g. have pretty Perl scraps, include the following L^AT_EX commands in your document:


```
\usepackage{listings}
...
\lstset{extendedchars=true, keepspaces=true, language=perl}
```

 See the `listings` documentation for a list of formatting options. Be sure to include a `\usepackage{listings}` in your document.
- `-V string` Provide *string* as the replacement for the `@v` operation.
- `-m` Emit a JSON map of the project to stdout and exit. The map describes every output file and every fragment, their definition locations, usage relationships, and transitive dependencies. Neither weaving nor tangling is performed. See Section 1.8 for details.
- `-e name` Extract a single fragment by name and emit its code and dependency information as Markdown to stdout. The match uses the same prefix logic as `@<name...@>`. Neither weaving nor tangling is performed. See Section 1.8 for details.
- `-R file[:line]` Read section markers from a tangled output file and emit JSON mapping to stdout. If a line number is given, emit the single innermost scrap covering that line. Without a line number, emit all scrap regions in the file. No `.weft` source files are needed—the reverse map operates entirely on the tangled output. See Section 1.8 for details.

1.5 Markdown Output

When invoked with `-w md`, `weft` generates a Markdown file (`.md`) instead of a L^AT_EX file (`.tex`). This allows literate programs to be read on any platform that renders Markdown, including GitHub, GitLab, and local editors like VS Code.

The Markdown output preserves all the essential features of the L^AT_EX output—scrap definitions, cross-references, and indices—but uses a different set of formatting conventions:

Scrap headers File definitions appear as `***filename*** **N** =` and fragment definitions as `***<Name* N>*** =`, where *N* is the sequential scrap number.

Scrap bodies Code is wrapped in `<pre><code>` HTML blocks (not fenced code blocks). This allows fragment references within scraps to be rendered as clickable HTML links with italic and bold formatting, which would not be possible inside Markdown fenced blocks.

HTML escaping Characters `<`, `>`, and `&` inside scraps are escaped as `<`, `>`, and `&` respectively.

Cross-references Each scrap has an HTML anchor (`<a id="weftN">`) and cross-reference links use `<a href="#weftN">`. Footer notes (“defined by”, “referenced in”) appear as Markdown blockquotes.

Indices The `@f`, `@m`, and `@u` commands produce Markdown lists with links. In the user identifier index, defining scraps appear in bold.

Sequential numbering The `-n` flag is forced automatically, since Markdown has no page numbers. Scraps are numbered 1, 2, 3, ... without letter suffixes.

Prose Text outside scraps passes through unmodified. The author writes Markdown prose directly in the `.weft` source file; `weft` does not interpret or transform it.

Syntax highlighting When the language of a scrap can be determined—from a preceding `@l` tag, the `@L` command, or the file extension of an `@o` name—`weft` emits `<code class="language-X">` on the opening tag. This allows client-side highlighters (`highlight.js`, `Prism`) and platforms like GitHub to apply syntax colouring automatically.

Tangling (output file generation) works identically regardless of the weave format. When `-w` is given, a single invocation produces both the documentation and the tangled output files in one pass.

1.6 Section Markers and Language Auto-Detection

When debugging tangled code, it is useful to correlate lines in the generated file back to the original web source. `weft` can auto-detect the programming language and generate *section markers*—comments that delimit the opening and closing of each scrap in the tangled output, including the source file name and line number. For C and C++, it also generates `#line` directives automatically.

The key idea is simple:

file extension → language → comment style → section markers in tangled output

For example, `@o server.js` has the extension `.js`. `weft` looks up `js` in its language table, finds that JavaScript uses `//` comments, and wraps each scrap in the tangled `server.js` with `//` markers:

```
// {1: web.weft:12}
console.log("hello");
// {1: }
```

Similarly, `@o main.c` matches `c`, which uses `/* */` comments and also enables `#line` directives:

```
#line 5 "core.weft"
/* {1: core.weft:5} */
#include <stdio.h>
/* {1: } */
```

And `@o config.py` matches `py`, using `#` comments:

```
# {3: settings.weft:40}
DATABASE_URL = "postgres://localhost/app"
# {3: }
```

This happens automatically—no flags are needed. The extension alone is enough.

1.6.1 Overriding the Language with `@l`

Sometimes the file extension is not enough, or you want a different behaviour. The `@l` command (lowercase) placed on the line before an `@o` or `@d` directive overrides the extension:

```
@l sql
@o queries.txt
@{SELECT * FROM users;
@}
```

Here the file is `queries.txt`, whose extension `.txt` has no comment style. But `@l sql` tells weft to use `--` comments, so the tangled output gets:

```
-- {1: db.weft:5}
SELECT * FROM users;
-- {:1}
```

`@l` also works on `@d` (fragment) definitions. This is useful for fragments that will be included in files of a specific language:

```
@l python
@d Validate email
@{import re
def validate(email):
    return bool(re.match(r'^[~@]+@[~@]+$', email))
@}
```

The language tag is stored on the fragment and used both for section markers in the tangled output and for syntax highlighting in the Markdown weave output.

`@l` is format-agnostic: it works identically in $\text{\LaTeX}$ and Markdown source files. Unlike the earlier approach of encoding the language in Markdown fences (`'''python`) or $\text{\LaTeX}$ markers (`%%python`), `@l` uses the same at-sign syntax as every other weft command.

Note that cosmetic fences (`'''`) may still be used around `@o/@d` definitions for readability in Markdown previewers—they are silently stripped during Markdown weaving (see Section 1.2.1). But fences no longer set the language; use `@l` for that.

Priority rules. When multiple sources of language information are available, weft resolves them in this order:

1. **Explicit per-file flags** (`-d`, `-s`, `-cx`, etc.) on the `@o` line always win for tangle behaviour.
2. **@l tag** immediately before the directive—overrides the extension.
3. **File extension** of the `@o` name—the default when no `@l` is present.

1.6.2 Mid-Line Inline Expansion

When a fragment reference `@<name>` appears in the middle of a line—after at least one visible character—weft automatically suppresses all tangle annotations for that expansion: section markers, `#line` directives, and fragment-name comments. The fragment content is spliced directly into the surrounding text, producing clean inline output.

This makes small constant fragments practical:

```
@d port @{4000@}
@d host @"localhost"@

@l c
@o config.c -s
@{const char *host = @<host>;
int port = @<port>;
@}
```

Tangling produces:

```
const char *host = "localhost";
int port = 4000;
```

If the same reference appears at the *start* of a line (even if indented), it is treated as a block expansion and markers are emitted normally. No special flags or directives are needed—the position on the line is the only thing that determines the behaviour.

This works with any language and any fragment size. When a multi-line fragment is expanded mid-line, its content is still spliced without markers—the rule is based on the *reference site*, not the fragment body.

1.6.3 The @L Command: Defining New Languages

The built-in table covers over 250 entries, but it cannot anticipate every language or every project’s preferences. The @L command defines new language mappings—or overrides existing ones—directly in the document preamble.

```
@L name style [+d]
```

- *name* is the tag that will be recognised in @l commands and file extensions (@o file.name).
- *style* is the comment syntax: //, #, --, /*, or <!--.
- +d (optional) enables #line directive generation.

Examples:

```
@L wgs1 //          % WebGPU shading language: // comments
@L cedar //         % Cedar policy language
@L c //             % override built-in: use // instead of /* */
```

How @L interacts with the rest of the system. User-defined entries are searched *before* the built-in table, so @L c // effectively changes the comment style for all subsequent @o main.c directives (and for any @l c tag) from /* */ to //. Up to 32 user-defined languages may be registered.

@L should appear in the document preamble, before any @o or @d that depends on it. It is processed during pass 1 and silently skipped during weaving.

1.6.4 The Language Table

weft includes a built-in table of over 250 entries (languages, file extensions, and common aliases), organised by comment style. Each entry maps a tag to a comment style and optionally enables #line directives. The tag is matched against @l tags and file extensions alike—there is no distinction between them.

Comment style	Tags (representative)
/* */ + #line	c, h, cpp, cc, cxx, hpp, hxx, objc, objcpp, m, mm, ino, pde, cu, cuh, cilk
/* */	css, scss, sass, less, stylus, cl (OpenCL)
//	js, ts, jsx, tsx, mjs, cjs, mts, svelte, vue, java, scala, kotlin, kt, groovy, gradle, clj, cs, fs, fsx, vb, go, rust, rs, zig, swift, dart, odin, mojo, carbon, php, graphql, prisma, glsl, hlsl, wgs1, metal, vert, frag, proto, thrift, fbs, capnp, sol, reason, res, ml, sml, fut, apex, bicep, jsonnet, qml, p4
#	python, py, pyx, pyi, sh, bash, zsh, fish, ksh, nu, perl, pl, ruby, rb, r, tcl, awk, sed, coffee, makefile, mk, cmake, just, meson, bazel, bzl, julia, jl, nim, crystal, cr, elixir, ex, yaml, yml, toml, dockerfile, terraform, tf, hcl, nix, conf, ini, env, editorconfig, gitignore, powershell, ps1, xdc, sdc
--	lua, luau, sql, psql, plsql, plpgsql, mysql, sqlite, hive, haskell, hs, elm, purescript, purs, ada, adb, ads, vhdl, vhd, lean, lean4, agda, idris, idr, erlang
<!-- -->	html, htm, xml, svg, xhtml, xsl, xsd, dtd, rss, xaml, csproj, jsp, aspx, cshtml, erb, ejs, hbs, njk, twig, blade, pug, liquid, mustache, md, mdx, rst, tex, latex
(none)	json, csv, tsv, txt, diff, patch

For the full list, see the `lang_table` array in Section ?? . If a tag is not found in either the user-defined (@L) or built-in table, weft defaults to `//` comments with no `#line` directives. Tags mapped to style 0 (e.g. `json`, `txt`) produce no markers at all.

1.6.5 Syntax Highlighting in Markdown Output

When weaving to Markdown (`-w md`), weft emits the language tag on the `<code>` element:

```
<pre><code class="language-python">def hello():
    print("hello")
</code></pre>
```

This allows client-side highlighters (highlight.js, Prism) and platforms like GitHub to apply syntax colouring automatically. The language tag comes from the same chain described above: @l tag or file extension.

For @d fragments, the tag comes from a preceding @l command (there is no file extension to infer from). If no language is known, weft emits a plain `<pre><code>` without a class.

1.6.6 Output Format of Section Markers

Each scrap in the tangled file is wrapped with an opening marker at the top and a closing marker at the bottom, using the comment syntax of the detected language. The opening marker includes the scrap number and the source location (file name and line number). When fragments are nested, each expansion gets its own pair of markers, producing a clear hierarchy in the output.

1.6.7 Debugging Tangled Code

When you work with literate programming, the code you execute is *tangled output*—generated files assembled from scraps scattered across `.w` files. An error at `server.js:47` does not correspond to line 47 of any file you wrote; it is line 47 of a file that weft assembled from fragments defined in `server.w`, `routes.w`, and `database.w`. Without markers, debugging means mentally reconstructing which scrap produced which line. With markers, the tangled file tells you directly.

Reading the markers. Each marker encodes three pieces of information:

Part	Meaning
{ or {:	Opening or closing delimiter
Number	Scrap ID—unique across the entire document
file:line	Source <code>.w</code> file and line number (opening only)

Examples by language. **JavaScript** (`//` comments):

```
// {3: routes.w:42}
app.get("/api/users", async (req, res) => {
  // {7: database.w:18}
  const users = await db.query("SELECT * FROM users");
  // {7}
  res.json(users);
});
// {3}
```

Python (`#` comments):

```
# {5: models.w:31}
class User:
    def __init__(self, name, email):
        self.name = name
```



```

        self.email = email
    # {:5}

```

C (`/* */` comments with `#line` directives):

```

#line 12 "memory.w"
/* {2: memory.w:12} */
void *safe_malloc(size_t size) {
    void *ptr = malloc(size);
    if (!ptr) {
        /* {4: errors.w:8} */
        fprintf(stderr, "out of memory\n");
        exit(1);
        /* {:4} */
    }
    return ptr;
}
/* {:2} */

```

Lua (`--` comments):

```

-- {9: config.w:55}
local defaults = {
    port = 8080,
    host = "localhost",
}
-- {:9}

```

HTML (`<!-- -->` comments):

```

<!-- {6: layout.w:20} -->
<nav class="sidebar">
    <!-- {11: navigation.w:7} -->
    <ul>
        <li><a href="/">Home</a></li>
    </ul>
    <!-- {:11} -->
</nav>
<!-- {:6} -->

```

Debugging workflow. The following four-step procedure traces any error in tangled output back to the `literate` source.

Step 1: Locate the error in tangled output. You receive an error such as:

```

TypeError: Cannot read property 'id' of undefined
    at Object.<anonymous> (server.js:47)

```

Open `server.js` and go to line 47.

Step 2: Read the nearest opening marker. Scan *upward* from the error line until you find an opening marker:

```

// {3: routes.w:42}

```

This tells you: scrap 3 generated this code, defined in `routes.w` starting at line 42.

Step 3: Go to the `literate` source. Open `routes.w` at line 42. You will find the scrap with the bug, surrounded by prose that explains its intent:

We fetch the user by ID from the database. If the user doesn't exist, we should return a 404, but currently we assume the user always exists --- this is the bug.

```
@d Handle user fetch
@{const user = await getUser(req.params.id);
  res.json(user.id);
@}
```

The prose makes the bug obvious: there is no null check. Fix it in the `.w` file, re-tangle, and the error is gone.

Step 4: Nested markers = nested fragments. When fragments invoke other fragments, markers nest:

```
// {3: routes.w:42}
app.get("/api/users/:id", async (req, res) => {
  // {7: database.w:18}
  const user = await db.query(
    "SELECT * FROM users WHERE id=$1",
    [req.params.id]);
  // {7}
  if (!user) return res.status(404).json({error: "not found"});
  res.json(user);
});
// {3}
```

If the bug is inside the `{7: ... {7}}` block, go to `database.w:18`, not `routes.w:42`. The *innermost enclosing markers* point to the most specific source location. Always look at the tightest pair around the error line.

Interpreting nested hierarchies. Complex tangled output can have several nesting levels. Read them like parentheses—each opening matches its corresponding close:

```
/* {1: main.w:10} */           outermost: main file structure
#include <stdio.h>
/* {2: io.w:25} */             middle: I/O subsystem
void print_report() {
  /* {5: format.w:8} */         innermost: formatting detail
  printf("Report v%d\n", VERSION);
  /* {5} */
}
/* {2} */
int main() {
  /* {3: init.w:14} */
  init_subsystems();
  /* {3} */
  print_report();
  return 0;
}
/* {1} */
```

Scrap numbers are *document-global*—scrap 1 is the first `@{...@}` in the entire woven document, scrap 5 is the fifth. They are stable as long as you do not reorder scraps.

C/C++: #line directives. For C-family languages, weft also generates `#line` directives *above* the opening marker. These make the C compiler itself report errors against the `.w` source file:

```
#line 12 "memory.w"
/* {2: memory.w:12} */
void *safe_malloc(size_t size) { ... }
/* {:2} */
```

Now gcc will say:

```
memory.w:14: error: implicit declaration of function 'malloc'
```

instead of:

```
memory.c:47: error: implicit declaration of function 'malloc'
```

The `#line` directive and the section marker are complementary: `#line` makes the compiler report errors against the `.w` file directly; the `{N: file:line}` markers are human-readable when browsing the tangled output.

Common debugging scenarios. Runtime error with stack trace.

```
Error: ENOENT: no such file or directory, open '/data/config.json'
    at readConfig (server.js:123)
```

Open `server.js:123`, find the nearest marker `// {12: config.w:35}`, open `config.w:35`—the scrap that reads the config—and fix the path or add an existence check in the `.w` file.

Compiler error in C. With `#line` directives active, the compiler already reports:

```
parser.w:84: error: 'token_type' undeclared
```

Go directly to `parser.w:84`. The prose around the scrap explains what `token_type` should be and where it is defined.

Wrong output, no error. The program runs but produces wrong results. Open the tangled file and search for the relevant code. The markers tell you which `.w` file and scrap to examine:

```
# {15: transform.w:67}
result = [x * 2 for x in data]    # should be x * 3?
# {:15}
```

Go to `transform.w:67`. The prose says “triple each value” but the code says `* 2`. Bug found.

Identifying which fragment owns a line. Sometimes you know the line of code but not which `@d` definition produced it (multiple fragments with the same name accumulate). The scrap number is unique:

```
// {14: handlers.w:89}
validateInput(data);
// {:14}
// {22: handlers.w:112}
validatePermissions(user);
// {:22}
```

Scraps 14 and 22 are two different definitions (or continuations) of the same fragment. The line numbers point to each one independently.

Tips.

- *Never edit tangled files.* The markers remind you these are generated. Always trace back to the `.w` source and edit there.
- *Scrap numbers are stable* as long as you do not add or remove `@{...@}` blocks before them. If numbering shifts after edits, re-tangle and the markers update automatically.
- *Use `grep` on markers* to find all scraps from a specific file:

```
grep -n '{[0-9]*: database.w:' server.js
```
- *Closing markers as verification.* If you see `{:7}` but no `{7:` above it in your viewport, scroll up—the scrap is larger than your screen. If you see `{7:` but no `{:7}`, the scrap extends below.
- *Markers in test output.* If your test framework shows code snippets on failure, the markers point straight to the literate source.
- *Disable markers selectively.* Use explicit flags on `@o` (e.g. `@o file -t`) with no comment flags to suppress markers for a specific output file when comments would interfere—for example, binary formats, data files, or files where comments are not valid syntax.

1.7 Restrictions

Because weft is intended to be a simple tool, a few restrictions have been established. Over time, some of these may be eliminated; others seem fundamental.

- The handling of errors is not completely ideal. In some cases, weft simply warns of a problem and continues; in other cases it halts immediately. This behavior should be regularized.
- weft warns about references to fragments that haven't been defined, but doesn't halt. The name of the fragment is included in the output file surrounded by `<>` signs. This seems most convenient for development, but may change in the future.
- File names and index entries should not contain any `@` signs.
- Fragment names may be (almost) any well-formed $\text{\TeX}$ string. It makes sense to change fonts or use math mode; however, care should be taken to ensure matching braces, brackets, and dollar signs.
- Anything is allowed in the body of a scrap; however, very long scraps (horizontally or vertically) may not typeset well.
- Temporary files (created for comparison to the eventual output files) are placed in the current directory. Since they may be renamed to an output file name, all the output files should be on the same file system as the current directory. Alternatively, you can use the `-p` flag to specify where the files go.
- Because page numbers cannot be determined until the document has been typeset, weft must be rerun after $\text{\LaTeX}$ to obtain a clean version of the document (very similar to the way it is sometimes necessary to rerun $\text{\LaTeX}$ to obtain an up-to-date table of contents after significant edits). weft will warn (in most cases) when this needs to be done; in the remaining cases, $\text{\LaTeX}$ will warn that labels may have changed. This restriction does not apply to Markdown output (`-w md`), which uses sequential numbering and requires no additional processing passes.

Very long scraps may be allowed to break across a page if declared with `@O` or `@D` (instead of `@o` and `@d`). This doesn't work very well as a default, since far too many short scraps will be broken across pages; however, as a user-controlled option, it seems very useful.

1.8 Project Introspection

The `-m` and `-e` flags support machine-readable introspection of a weft project. They run after pass 1 has built the name trees and do *not* invoke the weave or tangle phases.

1.8.1 Map (`-m`)

The command

```
weft -m project.weft
```

writes a JSON object to stdout with two arrays, `"files"` and `"fragments"`. For each entry the map records the name, language tag, definition scraps (with source file and line), and dependency relationships. File entries list the fragments they use; fragment entries list the scraps that invoke them and the fragments they reference in turn.

This gives an AI assistant or any external tool a complete dependency graph of the project in a single invocation, without having to parse the weft source files.

1.8.2 Extract (`-e`)

The command

```
weft -e "fragment name" project.weft
```

finds the named fragment (using prefix matching), then emits its scrap code and dependency information as Markdown to stdout. Referenced fragments are followed transitively, so the output contains the full closure of the requested fragment.

This allows an AI assistant to retrieve exactly the code it needs to understand or modify a specific fragment, together with all its dependencies, in one step.

1.8.3 Reverse Map (`-R`)

The `-m` and `-e` flags require the `.weft` source files. The reverse map provides the opposite direction: given a tangled output file, it parses the section markers already embedded in that file and reports which `.weft` source and line produced each region. No source files are needed and pass 1 is never invoked.

Single-line query.

```
weft -R server.js:47
```

emits a JSON object identifying the innermost scrap covering line 47:

```
{"scrap": 3, "file": "routes.weft", "line": 42}
```

If the line is outside any scrap region, an empty object `{}` is returned.

Full-file mapping.

```
weft -R server.js
```

emits a JSON object listing every scrap region in the file:

```
{"source": "server.js", "regions": [
  {"start": 2, "end": 15, "scrap": 3,
   "file": "routes.weft", "line": 42},
  {"start": 5, "end": 8, "scrap": 7,
   "file": "database.weft", "line": 18}
]}
```

Regions may overlap when fragments are nested.

This enables CI/CD pipelines to translate compiler errors back to the literate source automatically, and lets AI agents resolve tangled-output locations without loading the `.weft` files at all.

1.9 Acknowledgements

Several people have contributed their time, ideas, and debugging skills. In particular, the contributions of the following people are gratefully acknowledged: Osman Buyukisik, Manuel Carriba, Adrian Clarke, Tim Harvey, Michael Lewis, Walter Ravenek, Rob Shillingsburg, Kayvan Sylvan, Dominique de Waleffe, and Scott Warren. Of course, most of these people would never have heard of nuweb (the original tool) without the efforts of George Greenwade.

Since maintenance has been taken over by Marc Mengel, Simon Wright and Keith Harwood online contributions have been made by:

- Walter Brown <wb@fnal.gov>
- Nicky van Foreest <n.d.vanforeest@math.utwente.nl>
- Javier Goizueta <jgoizueta@jazzfree.com>
- Alan Karp <karp@hp.com>

Bibliography

- [1] Donald E. Knuth, *Literate programming*, The Computer Journal **27** (1984), no. 2, 97–111.
- [2] ———, *The T_EXbook*, Computers and Typesetting, vol. 1986aA, Addison-Wesley, Reading, MA, USA, 1986.
- [3] ———, *T_EX: The program*, Computers and Typesetting, vol. B, Addison-Wesley, Reading, MA, USA, 1986b1986.
- [4] ———, *METAFONT: THE PROGRAM*, Computers and Typesetting, vol. D, Addison-Wesley, Reading, MA, USA, 1986d1986.
- [5] Leslie Lamport, *L^AT_EXa document preparation system user's guide and reference manual*, Addison-Wesley, Reading, MA, USA, 1985.
- [6] Silvio Levy, *WEB adapted to C, another approach*, TUB **8** (1987), no. 1, 12–13.
- [7] Norman Ramsey, *Literate-programming tools need not be complex*, Report at `ftp.cs.princeton.edu` in `/reports/1991/351.ps.Z`. Software at `ftp.cs.princeton.edu` in `/pub/noweb.shar.Z` and at `bellcore.com` in `/pub/norman/noweb.shar.Z`. CS-TR-351-91, Department of Computer Science, Princeton University, August 1992, Submitted to *IEEE Software*.
- [8] Ross Williams, *Funnelweb user's manual*, `ftp.adelaide.edu.au` in `/pub/compression` and `/pub/funnelweb`, University of Adelaide, Adelaide, South Australia, Australia, 1992.